

# Char Drivers

# Char Drivers

## ❑ Major and Minor Numbers

- Major Number  
typically indicates the family of the device.
- Minor Number  
allows drivers to distinguish the various devices they manage.
- **dev\_t** holds device numbers, both the major and minor parts.  
a **32-bit** quantity with **12** bits for major number and **20** for the minor number.
- Dynamic allocation of Major Numbers  
already what used for **kernelnewbies.ko**,  
however there was a static way.
- Dynamic Allocation is preferred over static, as  
a randomly picked major number will lead to conflicts  
and trouble if this driver widely deployed.

```
root@karimeshapa-Inspiron-5537:/home/karimeshapa/ldd/kernelnewbies# ls -al /dev/kernelnewbies
crw----- 1 root root 511, 0 23:07 18 /dev/kernelnewbies
root@karimeshapa-Inspiron-5537:/home/karimeshapa/ldd/kernelnewbies#
```

### */include/linux/types.h*

```
13 typedef u32 __kernel_dev_t;
14
15 typedef __kernel_fd_set fd_set;
16 typedef __kernel_dev_t dev_t;
```

### */include/linux/kdev\_t.h*

```
7 #define MINORBITS 20
8 #define MINORMASK ((1U << MINORBITS) - 1)
9
10 #define MAJOR(dev) ((unsigned int) ((dev) >> MINORBITS))
11 #define MINOR(dev) ((unsigned int) ((dev) & MINORMASK))
12 #define MKDEV(ma,mi) (((ma) << MINORBITS) | (mi))
```

# Char Drivers

## ❑ Major and Minor Numbers, *Cont'd*

- The kernel uses ***chrdevs*** global variable to manage device number's allocation.
- ***alloc\_chrdev\_region()*** iterates ***chrdevs*** from last and find an empty entry to return as the major number.

## ❑ Char Device Registration

- Allocation APIs
  - alloc\_chrdev\_region()*** : register a range of char device numbers.
  - class\_create()*** : create a struct class to your device.
  - device\_create()*** : creates a device and registers it with sysfs.
  - cdev\_init()*** : initialize a cdev structure with fileops, making it ready to add to the system with ***cdev\_add()***.
  - cdev\_add()*** : add a char device to the system.
- Remove APIs
  - cdev\_del()***, ***device\_destroy()***, ***class\_destroy()***, ***unregister\_chrdev\_region()***

*/fs/char\_dev.c*

```
34 static struct char_device_struct {
35     struct char_device_struct *next;
36     unsigned int major;
37     unsigned int baseminor;
38     int minorct;
39     char name[64];
40     struct cdev *cdev; → → → /* will die */
41 } *chrdevs[CHRDEV_MAJOR_HASH_SIZE];
```

# Char Drivers

## ❑ Advanced Operations

- Ioctl  
user-space needs to control the device  
i.e. eject media, change a baud rate, report error information,...  
**ioctl** numbers can be any numbers however it's better  
to follow the kernel convention ***/include/uapi/asm-generic/ioctl.h***
- Sleeping  
sleeping process is accomplished through a data structure called a **wait queue**.  
***wait\_queue\_head\_t my\_queue;***  
***init\_waitqueue\_head(&my\_queue);***  
***wait\_event(queue, condition);***  
***wait\_event\_interruptible(queue, condition);***  
***wait\_event\_timeout(queue, condition, timeout);***  
***wait\_event\_interruptible\_timeout(queue, condition, timeout);***  
***void wake\_up(wait\_queue\_head\_t \*queue);***  
***void wake\_up\_interruptible(wait\_queue\_head\_t \*queue);***

*user-space*  
*sys/ioctl.h*



```
int ioctl(int fd, unsigned long cmd, ...);
```

# Char Drivers

## ❑ Advanced Operations, *Cont'd*

- Poll and select  
block a process until any of a given set of file descriptors becomes available for reading or writing.

***unsigned int (\*poll) (struct file \*filp, poll\_table \*wait);***

The driver adds a wait queue to the ***poll\_table*** structure by calling the function ***poll\_wait()***

```
static unsigned int chardrvs_poll(struct file *filp, poll_table *wait)
{
    struct chardrvs_dev *dev = (struct chardrvs_dev *)filp->private_data;
    unsigned int mask = 0;

    mutex_lock(&dev->lock);
    poll_wait(filp, &dev->wq_f, wait);
    mutex_unlock(&dev->lock);
    if (dev->avail)
        mask |= POLLOUT;
    return mask;
}
```

# Char Drivers

## ❑ Sequence File Interface

- The **seq\_file** interface assumes that you are creating a virtual file that steps through a sequence of items that must be returned to user space.
- Over time, **/proc** methods have become notorious for buggy implementations when the amount of output grows large.
- As a way of cleaning up the **/proc** code and making life easier for kernel programmers, the **seq\_file** interface was added.
- The kernel implements major functionality to be used by driver writer in order to **open, read, seek, release,...** to deal with **/proc** created files.
- **Why sequence files?**  
Implementations of **proc** files had an awkward limitation of not being able to “print” more than a single **page** of output.

Sequence files solve this problem by generating the “output” of a **proc** file as a sequence of writes, each of which can be up to a page in size, with no limit on the number of writes.

```
static int mydrv_open(struct inode *inode, struct file *file)
{
    return single_open(file, mydrv_show, NULL);
}

static struct proc_ops myproc_fops = {
    .proc_open = mydrv_open,
    .proc_read = seq_read,
    .proc_lseek = seq_lseek,
    .proc_release = single_release,
};
```