

Time Management

Have a look

- Before getting into **Time Management**, please take a look at
 - LK_Bird's Eye View session,

S7 Process Scheduler, Time, IPC.

Time Management

- General Timing Hardware
 - Real-time Clock (**RTC**)
Keeping track of the current time, date and use it as a timestamp for various resources in the system.
 - Timestamp Counter (**TSC**)
Implemented in every x86 microprocessor by means of a 64-bit register called **TSC**.

It counts the number of clock signals arriving on the CLK pin of the processor as the fact that one processor clock's frequency might not be the same as others makes it vary across processors.

native_calibrate_tsc() : CPU clock frequency is calculated during system boot.

/arch/x86/kernel/tsc.c

```
/**
 * native_calibrate_tsc
 * Determine TSC frequency via CPUID, else return 0.
 */
unsigned long native_calibrate_tsc(void)
{
    ...
}
```

Time Management

- General Timing Hardware
 - Programmable Interrupt Timer (**PIT**)
tasks that need to be carried out by the kernel at regular intervals.

In most machines the PIT keeps on issuing timer interrupts on IRQ0 periodically at approximately **100 Hz** frequency (**tick rate**) leaving more time for user-mode code (programs) to execute without interruption.

- High-precision event timer (**HPET**)
The HPET works with clock signals in excess of 10 Mhz, issuing interrupts once every **100 ns**.

Time Management

- Hardware abstraction
 - Every system has at least one clock counter which is managed by a kernel structure.
 - HW abstraction is provided by *struct clocksource*, this is the structure used for system time.
 - This structure provides callbacks to access and control this clock source *read()*, *enable()*, *disable()*, *suspend*, *and resume* routines.

/include/linux/clocksource.h

```
struct clocksource {
    u64      (*read)(struct clocksource *cs);
    u64      mask;
    u32      mult;
    u32      shift;
    u64      max_idle_ns;
    u32      maxadj;
    u32      uncertainty_margin;
#ifdef CONFIG_ARCH_CLOCKSOURCE_DATA
    struct arch_clocksource_data archdata;
#endif
    u64      max_cycles;
    const char *name;
    struct list_head list;
    int      rating;
    enum clocksource_ids id;
    enum vdso_clock_mode vdso_clock_mode;
    unsigned long flags;

    int      (*enable)(struct clocksource *cs);
    void      (*disable)(struct clocksource *cs);
    void      (*suspend)(struct clocksource *cs);
    void      (*resume)(struct clocksource *cs);
    void      (*mark_unstable)(struct clocksource *cs);
    void      (*tick_stable)(struct clocksource *cs);

    /* private: */
#ifdef CONFIG_CLOCKSOURCE_WATCHDOG
    /* Watchdog related data, used by the framework */
    struct list_head wd_list;
    u64      cs_last;
    u64      wd_last;
#endif
    struct module *owner;
};
```

Time Management

- Linux Timekeeping

- jiffies**

32-bit variable holds the number of ticks elapsed since system bootup.

jiffies_64 is used instead, which allows for thousands of millions of years before the overflow occurs.

Some APIs :

get_jiffies_64() : get the current value of jiffies.

time_after(), **time_before()**, **time_after_eq()**,
time_before_eq() : time comparison taking into account the possibility of wraparound.

jiffies_to_msecs(), **jiffies_to_usecs()**, **jiffies_to_nsecs()** :
convert jiffies to other time units such as ms, us, ns.

/include/linux/jiffies.h

/kernel/time/jiffies.c

```
u64 get_jiffies_64(void)
{
    unsigned int seq;
    u64 ret;

    do {
        seq = read_seqcount_begin(&jiffies_seq);
        ret = jiffies_64;
    } while (read_seqcount_retry(&jiffies_seq, seq));
    return ret;
}
```

/kernel/time/time.c

```
#define time_after(a,b) \
    (typecheck(unsigned long, a) && \
     typecheck(unsigned long, b) && \
     ((long)((b) - (a)) < 0))
#define time_before(a,b)    time_after(b,a)

#define time_after_eq(a,b) \
    (typecheck(unsigned long, a) && \
     typecheck(unsigned long, b) && \
     ((long)((a) - (b)) >= 0))
#define time_before_eq(a,b) time_after_eq(b,a)
```

```
unsigned int jiffies_to_msecs(const unsigned long j)
{
    #if HZ <= MSEC_PER_SEC && !(MSEC_PER_SEC % HZ)
        return (MSEC_PER_SEC / HZ) * j;
    #elif HZ > MSEC_PER_SEC && !(HZ % MSEC_PER_SEC)
        return (j + (HZ / MSEC_PER_SEC - 1)/(HZ / MSEC_PER_SEC));
    #else
        # if BITS_PER_LONG == 32
            return (HZ_TO_MSEC_MUL32 * j + (1ULL << HZ_TO_MSEC_SHR32) - 1) >>
                HZ_TO_MSEC_SHR32;
        # else
            return DIV_ROUND_UP(j * HZ_TO_MSEC_NUM, HZ_TO_MSEC_DEN);
        # endif
    #endif
}
```

Time Management

- Linux Timekeeping, *Cont'd*
 - Timeval and timespec
this current time is maintained by keeping the number of seconds elapsed since midnight of January 01, 1970 (called *epoch*).

The second elements in each of these represent the time elapsed since the last second in **us** and **ns** respectively *timespec*, *timeval*.

So, the time (counter value) read from the clock source needs to be accumulated and tracked.

The structures *tk_read_base*, *timekeeper* serve this purpose.

/include/uapi/linux/time.h

```
struct timespec {
    __kernel_old_time_t tv_sec;    /* seconds */
    long tv_nsec;                 /* nanoseconds */
};

struct timeval {
    __kernel_old_time_t tv_sec;    /* seconds */
    __kernel_suseconds_t tv_usec; /* microseconds */
};
```

/include/linux/timekeeper_internal.h

```
struct tk_read_base {
    struct clocksource *clock;
    u64 mask;
    u64 cycle_last;
    u32 mult;
    u32 shift;
    u64 xtime_nsec;
    ktime_t base;
    u64 base_real;
};

struct timekeeper {
    struct tk_read_base tkr_mono;
    struct tk_read_base tkr_raw;
    u64 xtime_sec;
    unsigned long ktime_sec;
    struct timespec64 wall_to_monotonic;
    ktime_t offs_real;
    ktime_t offs_boot;
    ktime_t offs_tai;
    s32 tai_offset;
    unsigned int clock_was_set_seq;
    u8 cs_was_changed_seq;
    ktime_t next_leap_ktime;
    u64 raw_sec;
    struct timespec64 monotonic_to_boot;
    ...
};
```

Time Management

- Timers
 - Timers sometimes called dynamic timers or kernel timers.
 - Timers are represented by struct ***timer_list***.
 - Steps using timers
 - 1) Define a timer
i.e. ***struct timer_list my_timer;***
 - 2) Fill out internal values (expires, callback)
i.e. ***my_timer.expires = jiffies + delay;***
i.e. ***my_timer.function = my_function;***
 - 3) Activate the timer
i.e. ***add_timer(&my_timer);***
 - Other APIs

mod_timer() : changes the expiration of a given timer

del_timer() : deactivate a timer prior to its expiration, not wait for handler execution on other CPUs.

del_timer_sync() : deactivate a timer and wait for the handler to finish.

include/linux/timer.h

```
struct timer_list {
    struct hlist_node entry;
    unsigned long expires;
    void (*function)(struct timer_list *);
    u32 flags;

#ifdef CONFIG_LOCKDEP
    struct lockdep_map lockdep_map;
#endif
};
```

/kernel/time/timer.c

```
int mod_timer(struct timer_list *timer, unsigned long expires)
{
    ...
}
EXPORT_SYMBOL(mod_timer);

int del_timer(struct timer_list *timer)
{
    ...
}
EXPORT_SYMBOL(del_timer);

int try_to_del_timer_sync(struct timer_list *timer)
{
    ...
}
EXPORT_SYMBOL(try_to_del_timer_sync);
```


Time Management

- Timers, *Cont'd*
 - Small Delays
void udelay(unsigned long usecs);
void ndelay(unsigned long nsecs);
void mdelay(unsigned long msecs);